

1 Customer Churn Analysis & Predictive Modeling for Syria Tel

1.1 Project Overview

Syria Tel is a telecommunication company- mobile network- provider in Syria, founded in 2000. Using the DataSet holding Syria Tel Customer Churn information available on Kaggle, this project aims to build a best in class Machine Learning Algorithm, that can predict which and how many customers will churn based on the information available on the dataset.

1.2 Business Understanding

As important to Telecommunications companies as it is to earn new customers to their services, is the need to retain current customers and preventing them from going to competitors. The act of leaving customers or loss of customers are referred to in this industry as "Churn", from a business perspective. The business would like to understand exactly what factors or situations contribute to churning, and most importantly, post identification of those factors, define precise strategic initiatives to retain customers. This project, therefore aims to help the business first in identifying the attributes and factors that cause the churn, and in turn, building models that can help the business predict it, so that in fact strategic business initiatives can be outlined for solution.

1.2.1 Business Modeling Objectives

1. Building a model than can predict who will churn with a high level of accuracy
2. Identifying important attributes or features that are key in predicting customer churn

1.2.2 Important Model Success Considerations

We are testing for churn. If a customer churns it means churn is positive. A negative is a no, the customer did not churn. What will be more important for this problem are a False Negatives, mening that our model labeled a customer as "not going to churn" when actually churned.

Recall therefore is an important metric in evaluating our model under False Negatives. Recall pertains to the rate at which the model makes correct predictions about customer churning. Aided by a C Matrix- the goal is to minimize false negatives. Failure to identify a customer who is about to churn is more costly from a business persepective, than wrongly classyging a non-churning customer. A good successful model therefore should have an 85% recall at the very least. Precision and accuracy are also metrics to have in mind.

2 EDA - Loading and Understanding the DataSet

```
In [1]: #install imblearn library
!pip install imbalanced-learn
```

```
Requirement already satisfied: imbalanced-learn in /Users/jonax/opt/anaconda3/envs/learn-env/lib/python3.8/site-packages (0.7.0)
Requirement already satisfied: scikit-learn>=0.23 in /Users/jonax/opt/anaconda3/envs/learn-env/lib/python3.8/site-packages (from imbalanced-learn) (0.23.2)
Requirement already satisfied: joblib>=0.11 in /Users/jonax/opt/anaconda3/envs/learn-env/lib/python3.8/site-packages (from imbalanced-learn) (0.17.0)
Requirement already satisfied: scipy>=0.19.1 in /Users/jonax/opt/anaconda3/envs/learn-env/lib/python3.8/site-packages (from imbalanced-learn) (1.5.2)
Requirement already satisfied: numpy>=1.13.3 in /Users/jonax/opt/anaconda3/envs/learn-env/lib/python3.8/site-packages (from imbalanced-learn) (1.18.5)
Requirement already satisfied: threadpoolctl>=2.0.0 in /Users/jonax/opt/anaconda3/envs/learn-env/lib/python3.8/site-packages (from scikit-learn>=0.23->imbalanced-learn) (2.1.0)
```

```
In [2]: #import libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.preprocessing import StandardScaler
from imblearn.over_sampling import SMOTE

from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, precision_score, recall_score
from sklearn.metrics import roc_curve, auc
from sklearn.metrics import classification_report, confusion_matrix, ConfusionMatrixDisplay
from sklearn.neighbors import KNeighborsClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.feature_selection import RFECV
from sklearn.pipeline import Pipeline

%matplotlib inline
```

```
In [3]: #loading DataSet first for preview and understanding
customer_df= pd.read_csv('Data/Customer_churn.csv')
customer_df.head()
```

Out[3]:

	state	account length	area code	phone number	international plan	voice mail plan	number vmail messages	total day minutes	total day calls	total day charge	...	tota ev call:
0	KS	128	415	382-4657	no	yes	25	265.1	110	45.07	...	9
1	OH	107	415	371-7191	no	yes	26	161.6	123	27.47	...	10
2	NJ	137	415	358-1921	no	no	0	243.4	114	41.38	...	11
3	OH	84	408	375-9999	yes	no	0	299.4	71	50.90	...	8
4	OK	75	415	330-6626	yes	no	0	166.7	113	28.34	...	12

5 rows × 21 columns



```
In [4]: #Understanding database shape
customer_df.shape
```

Out[4]: (3333, 21)

In [5]: *#Understanding database information on nulls and data types*

```
customer_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3333 entries, 0 to 3332
Data columns (total 21 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   state                                3333 non-null   object
1   account length                       3333 non-null   int64
2   area code                            3333 non-null   int64
3   phone number                         3333 non-null   object
4   international plan                   3333 non-null   object
5   voice mail plan                      3333 non-null   object
6   number vmail messages                3333 non-null   int64
7   total day minutes                    3333 non-null   float64
8   total day calls                      3333 non-null   int64
9   total day charge                     3333 non-null   float64
10  total eve minutes                    3333 non-null   float64
11  total eve calls                      3333 non-null   int64
12  total eve charge                     3333 non-null   float64
13  total night minutes                  3333 non-null   float64
14  total night calls                    3333 non-null   int64
15  total night charge                   3333 non-null   float64
16  total intl minutes                   3333 non-null   float64
17  total intl calls                     3333 non-null   int64
18  total intl charge                    3333 non-null   float64
19  customer service calls               3333 non-null   int64
20  churn                                3333 non-null   bool
dtypes: bool(1), float64(8), int64(8), object(4)
memory usage: 524.2+ KB
```

From the above we can conclude that of the 3333 rows and 21 columns , there are non-null values in each of the columns. We can change some value columns such as area code , from integer to object data type. And, we can check further.

In [6]: *# convert area code from integer to string*

```
customer_df['area code'] = customer_df['area code'].astype(object)
customer_df['area code'].dtype
```

Out[6]: dtype('O')

```
In [7]: #Checking for missing values  
customer_df.isna().sum()
```

```
Out[7]: state                                0  
account length                             0  
area code                                  0  
phone number                              0  
international plan                         0  
voice mail plan                           0  
number vmail messages                     0  
total day minutes                         0  
total day calls                           0  
total day charge                           0  
total eve minutes                         0  
total eve calls                           0  
total eve charge                           0  
total night minutes                       0  
total night calls                         0  
total night charge                        0  
total intl minutes                        0  
total intl calls                          0  
total intl charge                         0  
customer service calls                    0  
churn                                      0  
dtype: int64
```

From the above, this database does not have missing values.

```
In [8]: #checking for duplicate numbers using phone numbers  
customer_df.duplicated(subset='phone number').value_counts()
```

```
Out[8]: False      3333  
dtype: int64
```

From the above we can conclude the database has no duplications in phone numbers.


3 EDA- Analysis & Statistics

3.1 Churn Discovery

```
In [9]: # Discover churn and vizualize it
customer_df['churn'].value_counts()

#chart
fig, ax = plt.subplots(figsize=(10,5))
customer_df.groupby('churn').size().plot(kind='pie', autopct='%.2f', text=
ax.set_title('Churn Rate', fontsize=18)
```

```
Out[9]: Text(0.5, 1.0, 'Churn Rate')
```



We see that there is an imbalance in the churn rate, with approx 86% False and 14% True (they churn). Let's look at some stats and the distributions of statistics in the database

3.1.1 Call and churn discovery


We can look at the number of calls that make at which point customers will churn.

```
In [20]: churn_call = customer_df.groupby('customer service calls')['churn'].agg([
churn_call
```

Out[20]:

count	
customer service calls	
0	697
1	1181
2	759
3	429
4	166
5	66
6	22
7	9
8	2
9	2

```
In [24]: plt.figure(figsize=(10, 8))
splot = sns.barplot(x='customer service calls', y='churn',
                    data=customer_df, palette='rocket_r', ci=None)
# Add annotations to bars
for p in splot.patches:
    splot.annotate(format(p.get_height(), '.1f'),
                   (p.get_x() + p.get_width() / 2., p.get_height()),
                   ha = 'center', va = 'center',
                   xytext = (0, 9),
                   textcoords = 'offset points')
plt.title('Customer Service Calls & Churn')
plt.show()
```



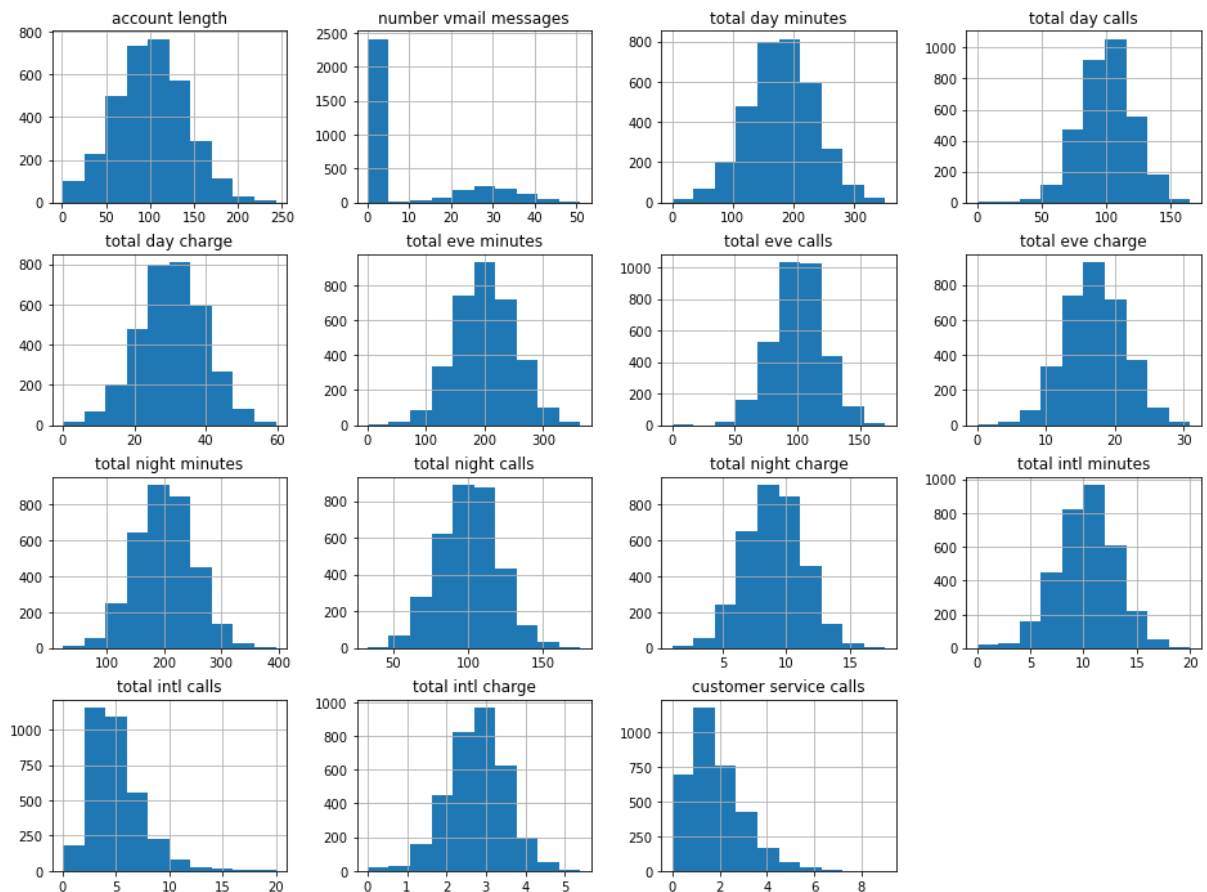
It takes about 4 calls, for customers to start churning, after that the probability of churning is above 50%


```
In [22]: # summary of statistics
customer_df.describe()
```

Out[22]:

	account length	number vmail messages	total day minutes	total day calls	total day charge	total eve minutes	total eve calls
count	3333.000000	3333.000000	3333.000000	3333.000000	3333.000000	3333.000000	3333.000000
mean	101.064806	8.099010	179.775098	100.435644	30.562307	200.980348	100.114311
std	39.822106	13.688365	54.467389	20.069084	9.259435	50.713844	19.922621
min	1.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
25%	74.000000	0.000000	143.700000	87.000000	24.430000	166.600000	87.000000
50%	101.000000	0.000000	179.400000	101.000000	30.500000	201.400000	100.000000
75%	127.000000	20.000000	216.400000	114.000000	36.790000	235.300000	114.000000
max	243.000000	51.000000	350.800000	165.000000	59.640000	363.700000	170.000000


```
In [26]: customer_df.drop(columns='churn').hist(figsize=(16,12));
```



While some distributions are normal approximately. Some are not and scaling and normalization are needed.

3.2 Multivariate Analysis

```
In [27]: #Heatmap to show correlations
fig, ax = plt.subplots(figsize=(12,12))
multi = np.triu(np.ones_like(customer_df.corr(), dtype=bool))
sns.heatmap(customer_df.corr(), linewidths=0.5, mask=multi, square=True,
```



Noting from the heatmap above four main views:

1) Perfect Correlations spotted: a) total day charge & total day minutes b) total even charge & total eve minutes c) total night charge & total nights minutes d) total international charge & total int. minutes One correlated variable from each of these pairs will have to be dropped from each pair to handle multicollinearity issues.

2) Churn - Weak Positive Correlation with total day minutes, total day charge, and customer service calls have a weak positive correlation with churn.

3) Churn- almost 0 correlations with most of features with exception of point 2 above.

4) Most features have very low correlations

4 Preparing Data for Machine Learning

4.0.1 Managing Multicollinearity

```
In [28]: # .75 < multicollinearity pairs
df = customer_df.corr().abs().stack().reset_index().sort_values(0, ascending=True)
df['pairs'] = list(zip(df.level_0, df.level_1))
df.set_index(['pairs'], inplace=True)
df.drop(columns=['level_1', 'level_0'], inplace=True)
df.columns = ['cc']
df.drop_duplicates(inplace=True)
df[(df.cc > .75) & (df.cc < 1)]
```

Out[28]:

	cc
pairs	
(total day charge, total day minutes)	1.000000
(total eve charge, total eve minutes)	1.000000
(total night minutes, total night charge)	0.999999
(total intl charge, total intl minutes)	0.999993

```
In [29]: # dropping a few columns to address collinearity
customer_df = customer_df.drop(columns = ['total day charge', 'total eve charge', 'total night charge', 'total intl charge'])
customer_df.columns
```

Out[29]: Index(['state', 'account length', 'area code', 'phone number', 'international plan', 'voice mail plan', 'number vmail messages', 'total day minutes', 'total day calls', 'total eve minutes', 'total eve calls', 'total night minutes', 'total night calls', 'total intl minutes', 'total intl calls', 'customer service calls', 'churn'], dtype='object')

4.0.2 Performing a train-test split

```
In [30]: #predictor and target variables
y= customer_df['churn']
X= customer_df.drop(columns= ['churn','phone number'])

#splitting data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25,
```

```
In [31]: #view into X_train head
X_train.head()
```

Out[31]:

	state	account length	area code	international plan	voice mail plan	number vmail messages	total day minutes	total day calls	total eve minutes	total eve calls	tot nig minuti
1066	KS	117	510	no	yes	25	216.0	140	224.1	69	267
1553	CO	86	415	no	no	0	217.8	93	214.7	95	228
2628	TN	37	415	no	no	0	221.0	126	204.5	110	118
882	FL	130	415	no	no	0	162.8	113	290.3	111	114
984	NV	77	415	no	no	0	142.3	112	306.3	111	196

```
In [32]: #Transforming categorical variables.Creating dummy variables for the categ
X_train= pd.get_dummies (X_train, drop_first= True)
X_test = pd.get_dummies (X_test, drop_first= True)
X_train.head(3)
```

Out[32]:

	account length	number vmail messages	total day minutes	total day calls	total eve minutes	total eve calls	total night minutes	total night calls	total intl minutes	total intl calls	...	state
1066	117	25	216.0	140	224.1	69	267.9	112	11.8	4	...	
1553	86	0	217.8	93	214.7	95	228.7	70	11.3	7	...	
2628	37	0	221.0	126	204.5	110	118.0	98	6.8	3	...	

3 rows × 65 columns

4.0.3 SMOTE


SMOTE is used to handle classes imbalance. We are trying to have a 50/50 split on into our training data(unlike what is seen in our Churn discovery entire data set of 86%False-14%True) . Therefore resampling X_train and y_training sets, and fitting SMOTE to these training data.

```
In [33]: #fitting SMOTE to training data.
smote= SMOTE(random_state=123)
X_train_resampled, y_train_resampled = smote.fit_resample(X_train, y_train)

#looking at the class distribution
print(pd.Series(y_train_resampled).value_counts())

#Visualize it
fig, ax = plt.subplots(figsize=(10, 5))
y_train_resampled.value_counts().plot(kind='pie', autopct='%0.2f', textprops=
ax.set_ylabel('Churn', fontsize=16)
ax.set_title('Churn distribution %', fontsize=20);
```

```
True      2127
False     2127
Name: churn, dtype: int64
```



In []:

Training sets churn distribution is balanced. Now we can ensure we don't have oversampled or undersampled. I can move forward to modeling the data.

5 Modeling Data

5.0.1 Metric Choice (Precision, Recall, Accuracy)

Choosing to measure Recall as the most appropriate for this kind of problem. False Negatives will cost more than False Positives. Not capturing someone who churned, losing them as a customer while paying a cost to acquire another customer to replace the lost customer is the highest challenge. If we identified someone as churned, when in fact they didn't and we do however pay to retain that False Positive would be less costly.

5.0.2 Baseline: Logistic Regression

Using 'liblinear', as solver.

```
In [34]: #creating pipeline
pipe_log = Pipeline(steps=[('scale', StandardScaler()), ('logreg', LogisticRegression(solver='liblinear'))])
pipe_log.fit(X_train_resampled, y_train_resampled)
```

```
Out[34]: Pipeline(steps=[('scale', StandardScaler()),
                          ('logreg',
                           LogisticRegression(fit_intercept=False, solver='liblinear'))])
```

5.0.2.1 Confusion Matrix and Model Evaluation


```
In [35]: #evaluating performance function & confusion matrix
def evaluate(model, X_test, y_test):
    y_train_preds = model.predict(X_train_resampled)
    y_test_preds = model.predict(X_test)
    print('Recall_score: ')
    print('Train: ', recall_score(y_train_resampled, y_train_preds))
    print('Test: ', recall_score(y_test, y_test_preds))
    print('\nPrecision_score: ')
    print('Train: ', precision_score(y_train_resampled, y_train_preds))
    print('Test: ', precision_score(y_test, y_test_preds))
    print('\nAccuracy_score: ')
    print('Train: ', accuracy_score(y_train_resampled, y_train_preds))
    print('Test: ', accuracy_score(y_test, y_test_preds))
    cm = confusion_matrix(y_test, y_test_preds, labels=model.classes_)
    disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=model.classes_)
    disp.plot();
```

```
In [36]: #passing the evaluation function
evaluate(pipe_log, X_test, y_test)
```

```
Recall_score:
Train:  0.8556652562294311
Test:   0.25225225225225223
```

```
Precision_score:
Train:  0.9494001043296818
Test:   0.50909090909090909
```

```
Accuracy_score:
Train:  0.9050305594734368
Test:   0.86810551558753
```



There is a large number of false negatives, making the model perform poorly. The model is overfitting performing well on the training data but not on the test data. On test data there is good accuracy at 86%, but recall is extremely low at 25%.

5.1 Decision Tree Model

5.1.1 Importing library and scaling data

```
In [37]: #Bringing Decision tree library from sklearn
from sklearn.tree import DecisionTreeClassifier
# Creating a pipeline
pipe_dt = Pipeline(steps=[('scale', StandardScaler()), ('clf', DecisionTr
pipe_dt.fit(X_train_resampled, y_train_resampled)
```

```
Out[37]: Pipeline(steps=[('scale', StandardScaler()),
                          ('clf',
                           DecisionTreeClassifier(criterion='entropy', random_sta
te=42))])
```

```
In [38]: #evaluate the performance of the model
```

```
evaluate(pipe_dt, X_test, y_test)
```

Recall_score:

Train: 1.0

Test: 0.7567567567567568

Precision_score:


Train: 1.0

Test: 0.49122807017543857

Accuracy_score:

Train: 1.0


Test: 0.8633093525179856



Clearly the decision tree has a higher performance than the baseline. It is still however overfitting.


```
In [39]: #Finding out the features that most affect the decision tree for Feature
def plot_feature_importances(pipe, figsize):
    model = pipe.steps[1][1]
    plt.figure(figsize=figsize)
    plt.barh(X_train_resampled.columns, model.feature_importances_, align
    plt.xlabel('Feature Importance')
    plt.ylabel('Feature')

plot_feature_importances(pipe_dt, (15, 15))
```



Top 3 features affecting churn include: 1)total day minutes 2) total mins and 3) customer service calls

5.1.2 Recursive Feature

This feature along with cross-validation is used to select a subset of features used to build model

```
In [40]: # feature selection using RFECV
rfecv = RFECV(estimator=DecisionTreeClassifier(random_state=42), scoring=
pipe_dt2 = Pipeline(steps=[('scale', StandardScaler()), ('Feature Selecti
pipe_dt2.fit(X_train_resampled, y_train_resampled)
```

```
Out[40]: Pipeline(steps=[('scale', StandardScaler()),
                          ('Feature Selection',
                           RFECV(estimator=DecisionTreeClassifier(random_state=4
2),
                                scoring='recall')),
                          ('clf', DecisionTreeClassifier(random_state=42))])
```

```
In [41]: # optimal number of features selected via cross-validation
print(f'Optimal # features: {rfecv.n_features_}' )
```

Optimal # features: 15

```
In [42]: # selected features (assigned a rank of 1)
rfecv_df = pd.DataFrame(rfecv.ranking_, index=X_train_resampled.columns, co
rfecv_df[rfecv_df['Rank'] == 1]
```

Out[42]:

	Rank
account length	1
area code_510	1
area code_415	1
international plan_yes	1
customer service calls	1
total intl calls	1
total intl minutes	1
voice mail plan_yes	1
total night minutes	1
total eve calls	1
total eve minutes	1
number vmail messages	1
total day minutes	1
total day calls	1
total night calls	1

```
In [43]: # remove columns that are not part of the optimal features
cols = rfecv_df[rfecv_df['Rank'] == 1].index
X_train_resampled = X_train_resampled[cols]
X_test = X_test[cols]
X_train_resampled.head(2)
```

Out[43]:

	account length	area code_510	area code_415	international plan_yes	customer service calls	total intl calls	total intl minutes	voice mail plan_yes	total night minutes	total eve calls
0	117	1	0	0	0	4	11.8	1	267.9	69
1	86	0	1	0	0	7	11.3	0	228.7	95

5.1.3 Hyperparameters testing in Decision Tree Model

In decision tree model the hyperparameters include max_depth, min_samples_leaf, min_sample_split, help the model be better tuned to stop the model from overfitting.

```
In [44]: #Using Gridsearch for Hyperparameters
# hyperparameter tuning using GridSearchCV
params_dt = {'clf__criterion': ['gini', 'entropy'],
             'clf__max_depth': range(14, 32, 2),
             'clf__min_samples_split': range(2, 10, 2),
             'clf__min_samples_leaf': [2, 3, 5, 7, 10],
             'clf__max_features': [11, 13, 15]}

gridsearch_dt = GridSearchCV(pipe_dt, params_dt, cv=4, scoring='recall')
gridsearch_dt.fit(X_train_resampled, y_train_resampled)
```

```
Out[44]: GridSearchCV(cv=4,
                    estimator=Pipeline(steps=[('scale', StandardScaler()),
                                              ('clf',
                                               DecisionTreeClassifier(criterion
n='entropy',
                                                                    random_s
tate=42))]),
                    param_grid={'clf__criterion': ['gini', 'entropy'],
                                'clf__max_depth': range(14, 32, 2),
                                'clf__max_features': [11, 13, 15],
                                'clf__min_samples_leaf': [2, 3, 5, 7, 10],
                                'clf__min_samples_split': range(2, 10, 2)},
                    scoring='recall')
```

```
In [45]: # parameters that gave the best result
print(f'Optimal parameters: {gridsearch_dt.best_params_}')

# Mean cross-validated score of the best_estimator
print(f'Validation recall: {gridsearch_dt.best_score_}')
```


```
Optimal parameters: {'clf__criterion': 'entropy', 'clf__max_depth': 28,
'clf__max_features': 15, 'clf__min_samples_leaf': 2, 'clf__min_samples_
split': 2}
Validation recall: 0.8725937371677781
```

```
In [46]: # evaluate the performance of the model
evaluate(gridsearch_dt, X_test, y_test)
```

```
Recall_score:
Train:  0.9811941701927598
Test:   0.7927927927927928
```

```
Precision_score:
Train:  0.9971333014811276
Test:   0.5238095238095238
```

```
Accuracy_score:
Train:  0.9891866478608369
Test:   0.8764988009592326
```



As opposed to the Decision Tree with no hyperparameters, this model has performed much better on recall score. The train and test scores are closer related, and overfitting has been reduced.

5.2 Random Forest Model

```
In [47]: #creating a pipeline
pipe_rf= Pipeline(steps=[('scale', StandardScaler()), ('rf', RandomForest
pipe_rf.fit(X_train_resampled, y_train_resampled)
```


```
Out[47]: Pipeline(steps=[('scale', StandardScaler()),
                          ('rf', RandomForestClassifier(random_state=42))])
```

```
In [48]: #evaluate model performance
evaluate(pipe_rf, X_test, y_test)
```

```
Recall_score:
Train:  1.0
Test:   0.6666666666666666
```

```
Precision_score:
Train:  1.0
Test:   0.6851851851851852
```

```
Accuracy_score:
Train:  1.0
Test:   0.9148681055155875
```



Random forest model has a recall score of 0.66. This is better than the baseline model of 0.25 recall on test data. However, in comparisson with Decision Tree Model with Hyperparameter recall of 0.79, it is not better.

5.2.1 Random Forest with Hyperparameters

```
In [49]: # hyperparameter tuning using GridSearchCV
params_rf = {'rf__n_estimators': range(400, 800, 200),
             'rf__criterion': ['gini', 'entropy'],
             'rf__max_depth': range(14, 20, 2),
             'rf__min_samples_split': range(3, 4, 7),
             'rf__min_samples_leaf': [5, 7, 12]}

gridsearch_rf = GridSearchCV(pipe_rf, params_rf, cv=4, scoring='recall')
gridsearch_rf.fit(X_train_resampled, y_train_resampled)
```

```
Out[49]: GridSearchCV(cv=4,
                    estimator=Pipeline(steps=[('scale', StandardScaler()),
                                              ('rf',
                                               RandomForestClassifier(random_s
tate=42))])),
        param_grid={'rf__criterion': ['gini', 'entropy'],
                    'rf__max_depth': range(14, 20, 2),
                    'rf__min_samples_leaf': [5, 7, 12],
                    'rf__min_samples_split': range(3, 4, 7),
                    'rf__n_estimators': range(400, 800, 200)},
        scoring='recall')
```

```
In [50]: # parameters that gave the best result
print(f'Optimal parameters: {gridsearch_rf.best_params_}')

# Mean cross-validated score of the best_estimator
print(f'Validation recall: {gridsearch_rf.best_score_}')


Optimal parameters: {'rf__criterion': 'entropy', 'rf__max_depth': 18,
                    'rf__min_samples_leaf': 5, 'rf__min_samples_split': 3, 'rf__n_estimator
s': 600}
Validation recall: 0.8702494229925095
```

```
In [51]: #evaluating the performance  
evaluate(gridsearch_rf, X_test, y_test)
```

```
Recall_score:  
Train:  0.9619181946403385  
Test:   0.6756756756756757
```

```
Precision_score:  
Train:  0.9898403483309144  
Test:   0.6696428571428571
```

```
Accuracy_score:  
Train:  0.9760225669957687  
Test:   0.9124700239808153
```



5.3 Model Finalization

The best model is the Decision Tree with Hyperparameters. This model has the highest recall scores in summary:

Recall scores: Test: 0.7927927927927928

Precision score: Test: 0.5238095238095238

Accuracy score: Test: 0.8764988009592326

Correctly identify a negative 80% of time. /100 , 80 of them identify and mis-identify 20%

6 Business Recommendations & Conclusions

According to our analysis, the most important features in predicting churn are :

- 1.The total number of minutes that the customer has done during the day

- 2.The total number of evening minutes
- 4.The number of customer calls to customer service
- 3.The total international minutes in international calls

For Syria tel, customer service is a key differentiator. Syria tel should focus on training programs to enhance this features that mostly need to reduce the number of minutes spent by a customer on the phone. Therefore effective communication training to customer service representatives is a key success factor. When a customer spends or makes 3 or more phone calls to customer service, this is an indication of higher churn. So providing customer incentives provided at the third call will be a recommendation.

6.0.1 Modeling- next steps recommendations & conclusions

The best performing model at a targeted level of 85% is not achieved by best model. There is still some overfitting. An increase in the training data set should reduce the overfitting and improve performance.